

Hyojung Han

AI Strategy & Platform Lead

hyojunguy@gmail.com · Jamsil-dong, Songpa-gu, Seoul, South Korea
github.com/hyojunguy · thakicloud.com/tech-blog · huggingface.co/ThakiCloud

SUMMARY

- Engineer of 19 unbroken years (2007–present) across models, infrastructure, and products. Now owns the strategy and architecture of three AI platforms at ThakiCloud (inference, training, agent automation) and leads the engineering team that builds them.
- Holds that the costliest mistake in AI decisions is trusting a wrong number: rebuilding the same recipe alone moves results 3.56pp, four FP4 toolchains silently compute in bf16, and a metric overlapping its synthetic data inflates the score. Catches these by measurement and hands the verdict to code gates, not to anyone's word.
- Treats AI unit cost as something measured, not estimated: serving configuration alone swings throughput 18.8x, and concurrency moves the cost of 1M output tokens from \$17.43 to \$4.59. Both were measured and turned back into product decisions.
- Founded a generative AI content platform, launched it across Japan, Taiwan, and Hong Kong, and grew it to 50K paying users and KRW 400M in revenue, carrying the technical call and the P&L on one desk. Inside the company, founded the internal AI tooling marketplace now shared by eight teams; outside it, 5 sole-authored arXiv papers, 34 open models, 1,328 blog posts.

PROFESSIONAL EXPERIENCE

ThakiCloud AI Platform Strategy, Architecture & Team Lead Seoul, South Korea Apr 2025 ~ Present

Owns strategy, architecture, and delivery across three AI platforms and leads the team that builds them. Writes the corporate strategy and the investor technical data room, and still trains and measures the models personally: 34 open releases and 4 papers.

Company AI strategy and team leadership

- Sole author of the corporate strategy document written for the executive team, defining the positioning of seven products and the vertical-integration thesis, then encoded that strategy as an always-on rule so every outward deliverable (blog, proposal, deck) passes through the same lens
- Leads a 7-person AI Platform team; argued single-point-of-failure risk per axis (serving, training, quantization, distillation, metering) into a headcount plan to 12 submitted to the executive team, and ran the hiring and onboarding schedule

Company-wide AI adoption: internal tooling platform and training

- Founded the internal AI tooling marketplace, packaging 150 skills and 14 agents for eight teams (sales, planning, frontend, backend, ML, QA, design, common); it is now maintained by 15 contributors as shared company infrastructure
- Designed and wrote the hands-on AI training curricula for three departments (11-13 modules each)

Korean 27B response-style alignment and its off-target effects (Human-KO)

- LoRA-SFT aligned Qwen3.8-27B for Korean response style: bullet usage 97.3% to 0.6%, median length 1,329 to 152 characters, human-likeness win rate 1.1% to 94.9% (n=175); the model only moved once the training set was cut from 2,992 records to 519 curated ones
- Found that style alignment shifts behaviors the objective never targeted, then corrected them with DPO: KoBBQ ambiguous-context abstention 81.5% to 92.8%, stereotyped answers 16.1% to 6.7%, with KMMLU up 4.2pp

Inference acceleration and FP4 kernel-path verification

- Added a block-diffusion speculative-decoding drafter (DFlash2) on top of my own NVFP4 quantization, lifting single-stream decode from 137.5 to 458.5 tok/s (3.33x), with the gain varying by prompt type (code 5.10x, reasoning 3.80x, chat 2.35x)
- Decomposed a prefix-cache fix and the drafter into three arms on a 234k-token context and showed the two effects are orthogonal: prefill 14.70s to 0.76s, decode 1.85x, end-user latency 15.32s to 1.09s

2i Studio (Founder) Founder Yongin, South Korea Mar 2024 ~ Apr 2025

Founded and ran a generative AI content platform end to end, from product to engineering to operations.

Built a face-based generative AI content platform

- Designed an identity-preserving generation pipeline on DreamBooth fine-tuning, growing to KRW 400M cumulative revenue and 50K paying users

Toss Senior Machine Learning Engineer Seoul, South Korea

Jan 2021 ~ Feb 2024

Designed and shipped machine learning systems across CDP personalization, generative AI tooling, face authentication, and workforce scheduling.

Customer data platform and personalization prediction

- Built per-service, per-user CTR/CVR prediction models, replacing manual SQL segment queries with automated probability scores

Samsung Electronics Senior Software Engineer, ML Researcher, Development Manager Suwon, South Korea

Sep 2011 ~ Dec 2020

Built an ML platform, 5G network automation, and a music recognition engine, while also shaping the organization's technical direction.

Built an internal cloud machine learning platform

- Designed a Kubernetes-based ML platform, cutting model deployment lead time from days to hours

Daum Machine Learning Engineer Jeju, South Korea

Sep 2008 ~ Aug 2011

Started in computer vision, building object recognition and image search systems.

Object recognition and image search system

- Designed a real-time image matching system at 7M-image scale using histogram-based Hamming embedding and LIS indexing

Vision Startup Software Engineer Seoul, South Korea

Jan 2007 ~ Aug 2008

Began my career in image processing and pattern recognition, the groundwork for later image search systems.

Image processing and pattern recognition modules

- Implemented C++ vision modules that extracted and matched features from camera input

PUBLICATIONS

Sole-authored, with reported numbers and reproduction code; four from ThakiCloud, one independent. Alongside them, 34 trained and compressed models on HuggingFace ThakiCloud: Human-KO, Satoori-KO, SKILLRET-Edge, vocabulary pruning, and serving quantization.

Off-Target Effects of Response-Style Alignment in a Korean 27B Language Model

arXiv:2609.11291 (Sep 2026)

On-Device Commercial Intent Retrieval Under Size, Latency, and Privacy Constraints: A 3 MiB Retrieval System with Typed Egress Boundaries

arXiv (submitted) · independent research

Rewrite Locality and Cache Provenance: Two Axes That Govern Prefix Reuse in Multi-Turn LLM Agents

arXiv (submitted)

Metric-Construction Coupling Inflates Measured Synthetic Dialect Recovery

arXiv (submitted)

Where Post-Training Quantization Breaks Text Embedders: A Measured Map Across Four Embedder Families

arXiv (submitted)

CORE TECHNICAL SKILLS

AI Strategy & Governance AI Unit-Cost Modeling / TCO · On-Prem vs Cloud Decisions · Vendor & Model Benchmark Design · Measurement-Integrity Gates · Safety & Bias Evaluation · Model & OSS License Review · Data Boundary & Privacy Design · Multi-Tenancy & Audit Trail

Training, Fine-tuning & Distillation SFT · CPT · DPO · GRPO · GKD · LoRA / QLoRA · Knowledge Distillation · FSDP / DDP

Inference, Serving & Compression vLLM · NVIDIA Dynamo · LMCache · Prefill-Decode Disaggregation · NVFP4 / W4A16 · LLM Compressor · MoE Compression / Pruning · Scale-to-Zero

Platform & GPU Infrastructure Kubernetes · Kueue / MultiKueue · ArgoCD / GitOps · Helm · Keycloak · SeaweedFS · CloudNativePG · NATS

EDUCATION

Yonsei University M.S. in Computer Science

2007